

13.5.1 Time Representation, Calendars, and Time Dimensions

For temporal databases, time is considered to be an *ordered sequence* of *points* in some *granularity* that is determined by the application. For example, suppose that some temporal application never requires time units that are less than one second. Then, each time point represents one second in time using this granularity. In reality, each second is a (short) *time duration*, not a point, since it may be further divided into milliseconds, microseconds, and so on. Temporal database researchers have used the term **chronon** instead of point to describe this minimal granularity for a particular application. The main consequence of choosing a minimal granularity—say, one second—is that events occurring within the same second will be considered to be *simultaneous events*, even though in reality they may not be.

Because there is no known beginning or ending of time, one needs a reference point from which to measure specific time points. Various calendars are used by various cultures (such as Gregorian (Western), Chinese, Islamic, Hindu, Jewish, Coptic, etc.) with different reference points. A **calendar** organizes time in different time units for convenience. Most calendars group 60 seconds into a minute, 60 minutes into an hour, and 24 hours into a day (based on the physical time of earth's rotation around its axis), and 7 days into a week. Further grouping of days into months and months into years either follow solar or lunar natural phenomena and are generally irregular. In the Gregorian calendar, which is used in most Western countries, days are grouped into months that are either 28, 29, 30, or 31 days, and 12 months are grouped into a year. Complex formulas are used to map the different time units to one another.

In SQL2, the temporal data types (see Chapter 7) include **DATE** (specifying Day, Month, and Year as DD-MM-YYYY), **TIME** (specifying Hour, Minute, and Second as HH:MM:SS), **TIMESTAMP** (specifying Date/Time combination, with options for including sub-second divisions if they are needed), **INTERVAL** (specifying relative time duration, such as 10 days or 250 minutes), and **PERIOD** (an *anchored* time duration with a fixed starting point, such as the 10-day period from January 1, 1999, to January 10, 1999, inclusive).⁸³

Event Information Versus Duration (or State) Information. A temporal database will store information concerning when certain events occur, or when certain facts are considered to be true. There are several different types of temporal information. **Point events** or **facts** are typically associated in the database with a **single time point** in some granularity. For example, a bank deposit event may be associated with the timestamp when the deposit was made, or the total monthly sales of a product (fact) may be associated with a particular month (say, February 1999). Note that even though such events or facts may have different granularities, each is still associated with a single time value in the database. This type of information is often represented as time series data as we shall discuss in Section 13.4.5. **Duration events** or **facts**, on the other hand, are associated with a specific time period in the database.⁸⁴ For example, an employee may have worked in a company from August 15, 1993, till November 20, 1998.

A time period is represented by its **start** and **end time points** [START-TIME, END-TIME]. For example, the above period is represented as [15-08-1993, 20-11-1998]. Such a time period is often interpreted to mean the *set of all time points* from start-time to end-time, inclusive, in the specified granularity. Hence, assuming a day granularity, the period [15-08-1993, 20-11-1998] represents the set of all days from August 15, 1993, until November 20, 1998, inclusive.⁸⁵

83. Unfortunately, the terminology has not been used consistently. For example, the term *interval* is often used to denote an anchored duration. For consistency, we shall use the SQL terminology.

84. This is the same as an anchored duration. It has also been frequently called a **time interval**, but to avoid confusion we will use **period** to be consistent with SQL terminology.

85. The representation [15-08-1993, 20-11-1998] is called a *closed interval* representation. One can also use an *open interval*, denoted [15-08-1993, 21-11-1998), where the set of points *does not include* the end point. Although the latter representation is sometimes more convenient, we shall use closed intervals throughout to avoid confusion.

Valid Time and Transaction Time Dimensions. Given a particular event or fact that is associated with a particular time point or time period in the database, the association may be interpreted to mean different things. The most natural interpretation is that the associated time is the time that the event occurred, or the period during which the fact was considered to be true in the real world. If this interpretation is used, the associated time is often referred to as the valid time. A temporal database using this interpretation is called a valid time database.

However, a different interpretation can be used, where the associated time refers to the time when the information was actually stored in the database; that is, it is the value of the system time clock when the information is valid in the system.⁸⁶ In this case, the associated time is called the transaction time. A temporal database using this interpretation is called a transaction time database.

Other interpretations can also be intended, but these two are considered to be the most common ones, and they are referred to as time dimensions. In some applications, only one of the dimensions is needed and in other cases both time dimensions are required, in which case the temporal database is called a bitemporal database. If other interpretations are intended for time, the user can define the semantics and program the applications appropriately, and it is called a user-defined time.

The next section shows with examples how these concepts can be incorporated into relational databases, and Section 13.4.3 shows an approach to incorporate temporal concepts into object databases.

15.2 Incorporating Time in Relational Databases Using Tuple Versioning

Valid Time Relations. Let us now see how the different types of temporal databases may be represented in the relational model. First, suppose that we would like to include the history of changes as they occur in the real world. Consider again the database in Figure 13.13, and let us assume that, for this application, granularity is day. Then, we could convert the two relations EMPLOYEE and DEPARTMENT into valid time relations by adding the attributes VST (Valid Start Time) and VET (Valid End Time), whose data type is DATE to provide day granularity. This is shown in Figure 13.19a, where the relations have been renamed EMP_VT and DEPT_VT, respectively.

Consider how the EMP_VT relation differs from the nontemporal EMPLOYEE relation (Figure 13.13).⁸⁷ In EMP_VT, each tuple v represents a version of an employee's information that is valid (in the real world) only for the time period [v.VST, v.VET], whereas in EMPLOYEE each tuple represents only the current state or current version of each employee. In EMP_VT, the current version of each employee typically has a special value, now, as its valid end time. This special value, now, is a temporal variable that implicitly represents the current time as time progresses. The nontemporal EMPLOYEE relation would only include those tuples from the EMP_VT relation whose VET is now.

Figure 13.20 shows a few tuple versions in the valid-time relations EMP_VT and DEPT_VT. There are two versions of Kumar, three versions of Chandra, one version of Agarwal, and one version of Narayan. We can see how a valid time relation should behave when information is changed. Whenever one or more tuples of an employee are updated, rather than actually overwriting the old values, as would happen in a nontemporal relation, the system should create a new version and close the current version by changing its valid end time. Hence, when the user issued the command to update the salary of Kumar effective on 2003-01-01 to Rs. 30000, the second version of Kumar was created (see Figure 13.20). At the time of this update, the first version of Kumar was the current version, with now as its VET, but after the update now was

more involved, as we shall see in Section 13.4.3. A nontemporal relation is also called a snapshot relation as it shows only the current snapshot or current state of the database.

changed to May 31, 2003 (one less than June 1, 2003, in day granularity), to indicate that the version has become a closed or history version and that the new (second) version of Kumar is now the current one.

(a) EMP_VT

NAME	<u>ENO</u>	SALARY	DNO	SUPERVISOR_ENO	<u>VST</u>	VET
------	------------	--------	-----	----------------	------------	-----

DEPT_VT

DNAME	<u>DNO</u>	TOTAL_SAL	MANAGER_ENO	<u>VST</u>	VET
-------	------------	-----------	-------------	------------	-----

(b) EMP_TT

NAME	<u>ENO</u>	SALARY	DNO	SUPERVISOR_ENO	<u>TST</u>	TET
------	------------	--------	-----	----------------	------------	-----

DEPT_TT

DNAME	<u>DNO</u>	TOTAL_SAL	MANAGER_ENO	<u>TST</u>	TET
-------	------------	-----------	-------------	------------	-----

(c) EMP_BT

NAME	<u>ENO</u>	SALARY	DNO	SUPERVISOR_ENO	VST	VET	<u>TST</u>	TET
------	------------	--------	-----	----------------	-----	-----	------------	-----

DEPT_BT

DNAME	<u>DNO</u>	TOTAL_SAL	MANAGER_ENO	VST	VET	<u>TST</u>	TET
-------	------------	-----------	-------------	-----	-----	------------	-----

FIGURE 13.19 Different types of temporal relational databases. (a) Valid time database schema. (b) Transaction time database schema. (c) Bitemporal database schema.

EMP_VT

NAME	<u>ENO</u>	SALARY	DNO	SUPERVISOR_ENO	<u>VST</u>	VET
Kumar	123456789	25000	5	333445555	15-06-2002	31-05-2003
Kumar	123456789	30000	5	333445555	01-06-2003	now
Chandra	333445555	25000	4	999887777	20-08-1999	31-01-2001
Chandra	333445555	30000	5	999887777	01-02-2001	31-03-2002
Chandra	333445555	40000	5	888665555	01-04-2002	now
Agarwal	222447777	28000	4	999887777	01-05-2001	10-08-2002
Narayan	666884444	38000	5	333445555	01-08-2003	now

...

DEPT_VT

DNAME	<u>DNO</u>	MANAGER_ENO	<u>VST</u>	VET
Research	5	888665555	20-09-2001	31-03-2002
Research	5	333445555	01-04-2002	now

...

FIGURE 13.20 Some tuple versions in the valid time relations EMP_VT and DEPT_VT.