

Experiment 8

Snashwat Shah

60004220126

TYBtech Comp B

Aim: To study and implement SparkSQL using PySpark.

Theory: SparkSQL is a query language for RDF data. It is used retrieve and manipulate stored data in RDF format.

Spark SQL stands for 'SparkSQL Protocol and RDF Query language'. It was developed by World Wide Web Consortium (W3C) and is a standard for querying RDF data. RDF stands for Resource description framework. It is a standard for describing resources on the web. RDF Data is stored in triplets which consist of a subject, a predicate and an object. SparkSQL query is used to retrieve data from an RDF dataset. It consists of a set of patterns are written in a syntax similar to SQL, but with some difference.

PySpark provides an easy-to-use interface to SparkSQL, allowing users to perform complex data processing tasks with few lines of code. With PySpark users can create Spark Dataframes, which is similar to Pandas Dataframes and can be queried using Spark SQL.

Conclusion: By studying and implementing Spark SQL using PySpark, we leverage the power of SQL queries to efficiently analyze and process big data within Spark enhancing data processing capabilities.

Experiment No.: 8

Aim: To study and Implement SparkQL using PySpark.

Code & Output:

IRIS DATASET:

```
[ ] from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("Online IRIS dataset").getOrCreate()

url = "/content/sample_data/Iris.csv"

df = spark.read.csv(url, header = False, inferSchema = True)

columns = ["id", "sepal_length", "sepal_width", "petal_length", "petal_width", "class"]
df = df.toDF(*columns)

df.createOrReplaceTempView("iris_data")

result1 = spark.sql("Select * from iris_data WHERE class = 'Iris-setosa'")
result2 = spark.sql("Select * from iris_data WHERE sepal_length > 7.0")
result3 = spark.sql("Select class, COUNT(*) from iris_data group by class ")

result1.show()
result2.show()
result3.show()

spark.stop()
```

```
+---+-----+-----+-----+-----+-----+
| id|sepal_length|sepal_width|petal_length|petal_width|      class|
+---+-----+-----+-----+-----+-----+
|  1|         5.1|         3.5|         1.4|         0.2|Iris-setosa|
|  2|         4.9|         3.0|         1.4|         0.2|Iris-setosa|
|  3|         4.7|         3.2|         1.3|         0.2|Iris-setosa|
|  4|         4.6|         3.1|         1.5|         0.2|Iris-setosa|
|  5|         5.0|         3.6|         1.4|         0.2|Iris-setosa|
|  6|         5.4|         3.9|         1.7|         0.4|Iris-setosa|
|  7|         4.6|         3.4|         1.4|         0.3|Iris-setosa|
|  8|         5.0|         3.4|         1.5|         0.2|Iris-setosa|
|  9|         4.4|         2.9|         1.4|         0.2|Iris-setosa|
| 10|         4.9|         3.1|         1.5|         0.1|Iris-setosa|
| 11|         5.4|         3.7|         1.5|         0.2|Iris-setosa|
| 12|         4.8|         3.4|         1.6|         0.2|Iris-setosa|
| 13|         4.8|         3.0|         1.4|         0.1|Iris-setosa|
| 14|         4.3|         3.0|         1.1|         0.1|Iris-setosa|
| 15|         5.8|         4.0|         1.2|         0.2|Iris-setosa|
| 16|         5.7|         4.4|         1.5|         0.4|Iris-setosa|
| 17|         5.4|         3.9|         1.3|         0.4|Iris-setosa|
| 18|         5.1|         3.5|         1.4|         0.3|Iris-setosa|
| 19|         5.7|         3.8|         1.7|         0.3|Iris-setosa|
| 20|         5.1|         3.8|         1.5|         0.3|Iris-setosa|
+---+-----+-----+-----+-----+-----+
only showing top 20 rows
```

```

+---+-----+-----+-----+-----+-----+
| id|sepal_length|sepal_width|petal_length|petal_width|      class|
+---+-----+-----+-----+-----+-----+
|103|      7.1|      3.0|      5.9|      2.1|Iris-virginica|
|106|      7.6|      3.0|      6.6|      2.1|Iris-virginica|
|108|      7.3|      2.9|      6.3|      1.8|Iris-virginica|
|110|      7.2|      3.6|      6.1|      2.5|Iris-virginica|
|118|      7.7|      3.8|      6.7|      2.2|Iris-virginica|
|119|      7.7|      2.6|      6.9|      2.3|Iris-virginica|
|123|      7.7|      2.8|      6.7|      2.0|Iris-virginica|
|126|      7.2|      3.2|      6.0|      1.8|Iris-virginica|
|130|      7.2|      3.0|      5.8|      1.6|Iris-virginica|
|131|      7.4|      2.8|      6.1|      1.9|Iris-virginica|
|132|      7.9|      3.8|      6.4|      2.0|Iris-virginica|
|136|      7.7|      3.0|      6.1|      2.3|Iris-virginica|
+---+-----+-----+-----+-----+-----+

+-----+-----+
|      class|count(1)|
+-----+-----+
|      Species|      1|
| Iris-virginica|      50|
|      Iris-setosa|      50|
|Iris-versicolor|      50|
+-----+-----+

```

Titanic Dataset - Perform SQL Queries to find:

- What is the number of passengers who survived the Titanic Disaster?
- How many female passengers were on board the Titanic?
- What is the average age of passengers in each passenger class?

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("Online IRIS dataset").getOrCreate()

url = "/content/sample_data/Titanic-Dataset.csv"

df = spark.read.csv(url, header = False, inferSchema = True)

columns = ["PassengerId", "Survived", "Pclass", "Name", "Sex", "Age", "SibSp", "Parch", "Ticket", "Fare", "Cabin", "Embarked"]
df = df.toDF(*columns)

df.createOrReplaceTempView("titanic")

result1 = spark.sql("Select * from titanic WHERE survived = 1")
result2 = spark.sql("Select count(*) from titanic where sex = 'female'")
result3 = spark.sql("Select avg(Age) from titanic group by Pclass ")

result1.show()
result2.show()
result3.show()

spark.stop()
```

PassengerId	Survived	Pclass	Name	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
2	1	1	Cumings, Mrs. Joh...	female	38	1	0	PC 17599	71.2833	C85	C
3	1	3	Heikkinen, Miss. ...	female	26	0	0	STON/O2. 3101282	7.925	NULL	S
4	1	1	Futrelle, Mrs. Ja...	female	35	1	0	113803	53.1	C123	S
9	1	3	Johnson, Mrs. Osc...	female	27	0	2	347742	11.1333	NULL	S
10	1	2	Nasser, Mrs. Nich...	female	14	1	0	237736	30.0708	NULL	C
11	1	3	Sandstrom, Miss. ...	female	4	1	1	PP 9549	16.7	G6	S
12	1	1	Bonnell, Miss. El...	female	58	0	0	113783	26.55	C103	S
16	1	2	Hewlett, Mrs. (Ma...	female	55	0	0	248706	16	NULL	S
18	1	2	Williams, Mr. Cha...	male	NULL	0	0	244373	13	NULL	S
20	1	3	Masselmani, Mrs. ...	female	NULL	0	0	2649	7.225	NULL	C
22	1	2	Beesley, Mr. Lawr...	male	34	0	0	248698	13	D56	S
23	1	3	McGowan, Miss. A...	female	15	0	0	330923	8.0292	NULL	Q
24	1	1	Sloper, Mr. Willi...	male	28	0	0	113788	35.5	A6	S
26	1	3	Asplund, Mrs. Car...	female	38	1	5	347077	31.3875	NULL	S
29	1	3	O'Dwyer, Miss. E...	female	NULL	0	0	330959	7.8792	NULL	Q
32	1	1	Spencer, Mrs. Wil...	female	NULL	1	0	PC 17569	146.5208	B78	C
33	1	3	Glynn, Miss. Mary...	female	NULL	0	0	335677	7.75	NULL	Q
37	1	3	Mamee, Mr. Hanna	male	NULL	0	0	2677	7.2292	NULL	C
40	1	3	Nicola-Yarred, Mi...	female	14	1	0	2651	11.2417	NULL	C
44	1	2	Laroche, Miss. Si...	female	3	1	2	SC/Paris 2123	41.5792	NULL	C

only showing top 20 rows

```
-----+
|count(1)|
-----+
|      314|
-----+
```

```
-----+
|      avg(Age)|
-----+
|      NULL|
| 25.14061971830986|
| 38.233440860215055|
| 29.87763005780347|
-----+
```

Wine Quality Dataset Example

- How many wines are considered high quality (quality score of 7 or higher)
- What is the average alcohol content of the wines in the dataset

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("Online IRIS dataset").getOrCreate()

url = "/content/sample_data/winequality-red.csv"

df = spark.read.csv(url, header = False, inferSchema = True)

columns = ["fixed acidity", "volatile acidity", "citric acid", "residual sugar", "chlorides", "free sulfur dioxide", "total sulfur dioxide", "density", "ph", "sulphates", "alcohol", "quality"]
df = df.toDF(*columns)

df.createOrReplaceTempView("wine")

result1 = spark.sql("Select count(*) from wine WHERE quality >= 7.0")
result3 = spark.sql("Select avg(alcohol) from wine ")

result1.show()
result3.show()

spark.stop()
```

```
+-----+
|count(1)|
+-----+
|      217|
+-----+

+-----+
|  avg(alcohol)|
+-----+
|10.422983114446502|
+-----+
```

California Housing Dataset Example:

- How many houses have a median value above \$5000,000 in California
- What is the average age of the houses in the dataset

```
from pyspark.sql import SparkSession

spark = SparkSession.builder.appName("Online IRIS dataset").getOrCreate()

url = "/content/sample_data/california_housing_train.csv"

df = spark.read.csv(url, header = False, inferSchema = True)

columns = ["longitude", "latitude", "housing_median_age", "total_rooms", "total_bedrooms", "population", "households", "median_income", "median_house_value"]
df = df.toDF(*columns)

df.createOrReplaceTempView("california")

result1 = spark.sql("Select count(*) from california WHERE median_house_value > 500000")
result3 = spark.sql("Select avg(housing_median_age) from california ")

result1.show()
result3.show()

spark.stop()
```

```
+-----+
|count(1)|
+-----+
|  16820|
+-----+

+-----+
|avg(housing_median_age)|
+-----+
|  28.58935294117647|
+-----+
```